

```

import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

# Reading the CSV file
student_per = pd.read_csv("/content/student-mat.csv", delimiter=";")
student_per

>{"type": "dataframe", "variable_name": "student_per"}

student_per.shape #This dataset has 395 samples and 33 features.

(395, 33)

student_per.info() #This gives the information about the columns of dataset, Using this we find all the non-numerical columns in dataset
#The given dataset doesn't have any null values.

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 395 entries, 0 to 394
Data columns (total 33 columns):
 #   Column      Non-Null Count  Dtype  
---  --  
 0   school      395 non-null    object  
 1   sex          395 non-null    object  
 2   age          395 non-null    int64  
 3   address     395 non-null    object  
 4   famsize     395 non-null    object  
 5   Pstatus      395 non-null    object  
 6   Medu         395 non-null    int64  
 7   Fedu         395 non-null    int64  
 8   Mjob          395 non-null    object  
 9   Fjob          395 non-null    object  
 10  reason        395 non-null    object  
 11  guardian     395 non-null    object  
 12  traveltime   395 non-null    int64  
 13  studytime    395 non-null    int64  
 14  failures      395 non-null    int64  
 15  schoolsup    395 non-null    object  
 16  famsup        395 non-null    object  
 17  paid          395 non-null    object  
 18  activities    395 non-null    object  
 19  nursery        395 non-null    object  
 20  higher         395 non-null    object  
 21  internet      395 non-null    object  
 22  romantic      395 non-null    object  
 23  famrel        395 non-null    int64  
 24  freetime       395 non-null    int64  
 25  goout          395 non-null    int64  
 26  Dalc           395 non-null    int64  
 27  Walc           395 non-null    int64

```

```

28  health      395 non-null    int64
29  absences    395 non-null    int64
30  G1          395 non-null    int64
31  G2          395 non-null    int64
32  G3          395 non-null    int64
dtypes: int64(16), object(17)
memory usage: 102.0+ KB

student_per.describe()

{"summary": {"name": "student_per", "rows": 8, "fields": [{"column": "age", "dtype": "number", "min": 134.436252896189, "max": 395.0, "num_unique_values": 8, "samples": [16.696202531645568, 17.0, 395.0], "semantic_type": "\\", "description": "\n"}, {"column": "Medu", "properties": {"dtype": "number", "std": 138.80963938157987, "min": 0.0, "max": 395.0, "num_unique_values": 7, "samples": [2.749367088607595, 3.0], "semantic_type": "\\", "description": "\n"}, {"column": "Fedu", "properties": {"dtype": "number", "std": 138.92085462409693, "min": 0.0, "max": 395.0, "num_unique_values": 7, "samples": [2.5215189873417723, 3.0], "semantic_type": "\\", "description": "\n"}, {"column": "traveltime", "properties": {"dtype": "number", "std": 139.0946757987501, "min": 0.6975047549086848, "max": 395.0, "num_unique_values": 6, "samples": [1.4481012658227848, 4.0], "semantic_type": "\\", "description": "\n"}, {"column": "studytime", "properties": {"dtype": "number", "std": 139.00700274471274, "min": 0.839240346418556, "max": 395.0, "num_unique_values": 6, "samples": [395.0, 4.0], "semantic_type": "\\", "description": "\n"}, {"column": "failures", "properties": {"dtype": "number", "std": 139.4513615014189, "min": 0.0, "max": 395.0, "num_unique_values": 5, "samples": [0.3341772151898734, 3.0], "semantic_type": "\\", "description": "\n"}, {"column": "famrel", "properties": {"dtype": "number", "std": 0.743650973606249, "min": 0.0, "max": 395.0, "num_unique_values": 5, "samples": [0.3341772151898734, 3.0], "semantic_type": "\\", "description": "\n"}], "samples": [395.0, 4.0], "semantic_type": "\\", "description": "\n"}]

```

```
138.45880901426744,\n      \"min\": 0.8966586076885056,\n      \"max\": 395.0,\n      \"samples\": [\n        395.0,\n        3.9443037974683546,\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\",\n      \"freetime\", \"properties\": {\n        \"number\":,\n        \"std\": 138.63828826062982,\n        \"min\": 0.9988620396657205,\n        \"max\": 395.0,\n        \"num_unique_values\": 7,\n        \"samples\": [\n          395.0,\n          4.0\n        ],\n        \"semantic_type\": \"\",,\n        \"description\": \"\",\n        \"column\": \"goout\", \"properties\": {\n          \"dtype\": \"number\",,\n          \"std\": 138.68948196584594,\n          \"min\": 1.0,\n          \"max\": 395.0,\n          \"num_unique_values\": 8,\n          \"samples\": [\n            395.0,\n            3.0\n          ],\n          \"semantic_type\": \"\",,\n          \"description\": \"\",\n          \"column\": \"Dalc\", \"properties\": {\n            \"dtype\": \"number\",,\n            \"std\": 139.0354623650101,\n            \"min\": 0.8907414280909659,\n            \"max\": 395.0,\n            \"num_unique_values\": 6,\n            \"samples\": [\n              395.0,\n              5.0\n            ],\n            \"semantic_type\": \"\",,\n            \"description\": \"\",\n            \"column\": \"Walc\", \"properties\": {\n              \"dtype\": \"number\",,\n              \"std\": 138.87302263653973,\n              \"min\": 1.0,\n              \"max\": 395.0,\n              \"num_unique_values\": 7,\n              \"samples\": [\n                395.0,\n                3.0\n              ],\n              \"semantic_type\": \"\",,\n              \"description\": \"\",\n              \"column\": \"health\", \"properties\": {\n                \"dtype\": \"number\",,\n                \"std\": 138.50262599778412,\n                \"min\": 1.0,\n                \"max\": 395.0,\n                \"num_unique_values\": 7,\n                \"samples\": [\n                  395.0,\n                  4.0\n                ],\n                \"semantic_type\": \"\",,\n                \"description\": \"\",\n                \"column\": \"absences\", \"properties\": {\n                  \"dtype\": \"number\",,\n                  \"std\": 136.85777166785417,\n                  \"min\": 0.0,\n                  \"max\": 395.0,\n                  \"num_unique_values\": 7,\n                  \"samples\": [\n                    395.0,\n                    8.0\n                  ],\n                  \"semantic_type\": \"\",,\n                  \"description\": \"\",\n                  \"column\": \"G1\", \"properties\": {\n                    \"dtype\": \"number\",,\n                    \"std\": 136.30663508587594,\n                    \"min\": 3.0,\n                    \"max\": 395.0,\n                    \"num_unique_values\": 8,\n                    \"samples\": [\n                      395.0,\n                      10.90886075949367,\n                      11.0,\n                      395.0\n                    ],\n                    \"semantic_type\": \"\",,\n                    \"description\": \"\",\n                    \"column\": \"G2\", \"properties\": {\n                      \"dtype\": \"number\",,\n                      \"std\": 136.4163745266465,\n                      \"min\": 0.0,\n                      \"max\":
```

```

395.0,\n          "num_unique_values": 8,\n          "samples": [\n10.713924050632912,\n           11.0,\n           395.0\n        ],\n        "semantic_type": "\",\n          "description": \"\"\n      }\n    },\n    {\n      "column": "G3",\n      "properties": {\n        "dtype": "number",\n        "std": 136.35024783099098,\n        "min": 0.0,\n        "max": 395.0,\n        "num_unique_values": 8,\n        "samples": [\n          10.415189873417722,\n          11.0,\n          395.0\n        ],\n        "semantic_type": "\",\n        "description": \"\"\n      }\n    }\n  ],\n  "type": "dataframe"
}

# Converting all the non-numerical columns in dataset to numerical
# columns using Label Encoder
from sklearn.preprocessing import LabelEncoder
label = LabelEncoder()
encode_colm = student_per.iloc[:, [0,1,3,4,5,8,9,10,11,15,16,17,18,19,20,21,22]]
for i in encode_colm:
    student_per[i] = label.fit_transform(student_per[i])
student_per

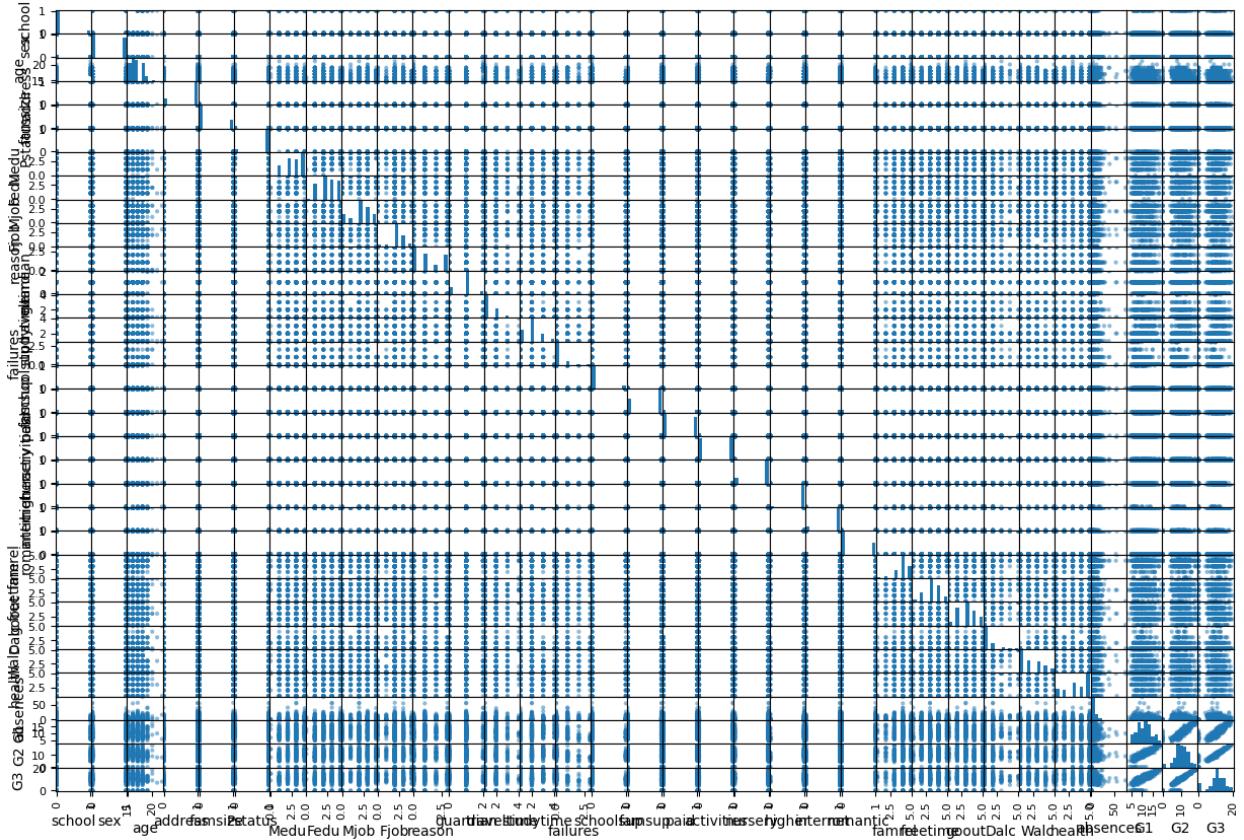
{"type": "dataframe", "variable_name": "student_per"}

# Correlation Matrix
correlation_Matrix = student_per.corr()
correlation_Matrix

{"type": "dataframe", "variable_name": "correlation_Matrix"}

# Scatter Matrix
pd.plotting.scatter_matrix(student_per, figsize=(15, 10))
plt.show()

```



```

# Storing the input attributes and the target variable
X = student_per.iloc[:, :32] # Input attributes
y = student_per.iloc[:, 32] # Target variable(G3)
X

{"type": "dataframe", "variable_name": "X"}

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test =
train_test_split(X, y, test_size=0.20, random_state=42)

from sklearn.preprocessing import StandardScaler
stand_scale = StandardScaler()
stand_scale.fit(X_train)
X_train_std = stand_scale.transform(X_train)
X_test_std = stand_scale.transform(X_test)

from sklearn.linear_model import LinearRegression
LR_model = LinearRegression()
LR_model.fit(X_train_std, y_train)
LR_score = LR_model.score(X_test_std, y_test) # Finding the R2 score of
# the model
print("R2 score:", LR_score)

R2 score: 0.7545777855043497

```

SUPPORT VECTOR MACHINES

```
from sklearn.svm import SVR
# SVM using rbf kernel
SVR_rbf = SVR(kernel='rbf')
# SVM using linear kernel
SVR_lin = SVR(kernel='linear')
# SVM using polynomial kernel
SVR_poly = SVR(kernel='poly')

SVR_rbf.fit(X_train_std,y_train)
SVR_lin.fit(X_train_std,y_train)
SVR_poly.fit(X_train_std,y_train)

# Finding the R2 score of SVR model using rbf kernel
rbf_score = SVR_rbf.score(X_test_std,y_test)
# Finding the R2 score of SVR model using linear kernel
lin_score = SVR_lin.score(X_test_std,y_test)
# Finding the R2 score of SVR model using polynomial kernel
poly_score = SVR_poly.score(X_test_std,y_test)

print("R2 score of SVM (rbf):",rbf_score)
print("R2 score of SVM (linear):",lin_score)
print("R2 score of SVM (polynomial):",poly_score)

R2 score of SVM (rbf): 0.6681792330768994
R2 score of SVM (linear): 0.7899692272521163
R2 score of SVM (polynomial): 0.49708883959770733
```

DECISION TREES

```
from sklearn.tree import DecisionTreeRegressor
tree_mod = DecisionTreeRegressor() # Performing DecisionTree Regressor
model
tree_mod.fit(X_train,y_train)
tree_score = tree_mod.score(X_test,y_test) # Finding the R2 score
print("R2 score:",tree_score)

R2 score: 0.6845481824149033

# Visualising the Regression tree
feature_names = student_per.columns[:32]
from sklearn.tree import export_graphviz
regression_tree = export_graphviz(tree_mod,out_file="data\studentperformance.dot",feature_names=feature_names,rounded=True,filled=True)

from graphviz import Source
Source.from_file("data\studentperformance.dot")
```

