

```

import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

# Reading the CSV file
student_per = pd.read_csv("/content/student-mat.csv", delimiter=";")
student_per

>{"type": "dataframe", "variable_name": "student_per"}

student_per.shape #This dataset has 395 samples and 33 features.

(395, 33)

student_per.info() #This gives the information about the columns of dataset, Using this we find all the non-numerical columns in dataset
#The given dataset doesn't have any null values.

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 395 entries, 0 to 394
Data columns (total 33 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          --          --      
 0   school      395 non-null    object  
 1   sex          395 non-null    object  
 2   age          395 non-null    int64  
 3   address     395 non-null    object  
 4   famsize     395 non-null    object  
 5   Pstatus      395 non-null    object  
 6   Medu         395 non-null    int64  
 7   Fedu         395 non-null    int64  
 8   Mjob          395 non-null    object  
 9   Fjob          395 non-null    object  
 10  reason        395 non-null    object  
 11  guardian     395 non-null    object  
 12  traveltime   395 non-null    int64  
 13  studytime    395 non-null    int64  
 14  failures      395 non-null    int64  
 15  schoolsup    395 non-null    object  
 16  famsup        395 non-null    object  
 17  paid          395 non-null    object  
 18  activities    395 non-null    object  
 19  nursery        395 non-null    object  
 20  higher         395 non-null    object  
 21  internet      395 non-null    object  
 22  romantic      395 non-null    object  
 23  famrel        395 non-null    int64  
 24  freetime       395 non-null    int64  
 25  goout          395 non-null    int64  
 26  Dalc           395 non-null    int64  
 27  Walc           395 non-null    int64

```

```

28  health      395 non-null    int64
29  absences    395 non-null    int64
30  G1          395 non-null    int64
31  G2          395 non-null    int64
32  G3          395 non-null    int64
dtypes: int64(16), object(17)
memory usage: 102.0+ KB

student_per.describe()

{"summary": {"name": "student_per", "rows": 8, "fields": [{"column": "age", "dtype": "number", "min": 134.436252896189, "max": 395.0, "num_unique_values": 8, "samples": [16.696202531645568, 17.0, 395.0], "semantic_type": "\\", "description": "\n"}, {"column": "Medu", "properties": {"dtype": "number", "std": 138.80963938157987, "min": 0.0, "max": 395.0, "num_unique_values": 7, "samples": [2.749367088607595, 3.0], "semantic_type": "\\", "description": "\n"}, {"column": "Fedu", "properties": {"dtype": "number", "std": 138.92085462409693, "min": 0.0, "max": 395.0, "num_unique_values": 7, "samples": [2.5215189873417723, 3.0], "semantic_type": "\\", "description": "\n"}, {"column": "traveltime", "properties": {"dtype": "number", "std": 139.0946757987501, "min": 0.6975047549086848, "max": 395.0, "num_unique_values": 6, "samples": [1.4481012658227848, 4.0], "semantic_type": "\\", "description": "\n"}, {"column": "studytime", "properties": {"dtype": "number", "std": 139.00700274471274, "min": 0.839240346418556, "max": 395.0, "num_unique_values": 6, "samples": [395.0, 4.0], "semantic_type": "\\", "description": "\n"}, {"column": "failures", "properties": {"dtype": "number", "std": 139.4513615014189, "min": 0.0, "max": 395.0, "num_unique_values": 5, "samples": [0.3341772151898734, 3.0], "semantic_type": "\\", "description": "\n"}, {"column": "famrel", "properties": {"dtype": "number", "std": 0.743650973606249, "min": 0.0, "max": 395.0, "num_unique_values": 5, "samples": [0.3341772151898734, 3.0], "semantic_type": "\\", "description": "\n"}], "samples": [395.0, 4.0], "semantic_type": "\\", "description": "\n"}]

```

```

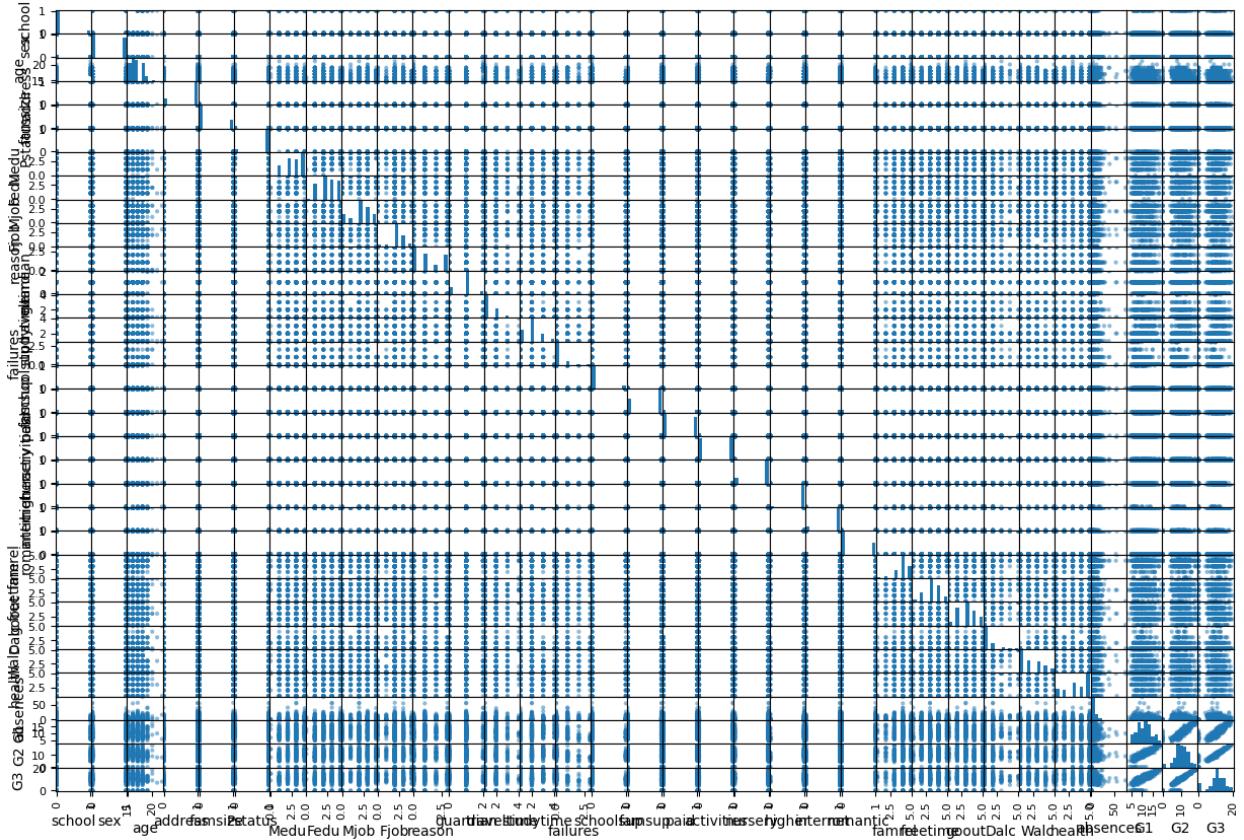
138.45880901426744, \n          \\"min\\": 0.8966586076885056, \n
\\"max\\": 395.0, \n          \\"num_unique_values\\": 6, \n
\\"samples\\": [\n            395.0, \n            3.9443037974683546, \n
5.0\n        ], \n          \\"semantic_type\\": \"\", \n
\\"description\\": \"\\n      \"}, \n      {\n        \\"column\\": \n\"
\\\"freetime\\\", \n        \\"properties\\": {\n          \\"dtype\\": \n\"
\\\"number\\\", \n          \\"std\\": 138.63828826062982, \n          \\"min\\": \n\"
0.9988620396657205, \n          \\"max\\": 395.0, \n
\\"num_unique_values\\": 7, \n          \\"samples\\": [\n            395.0, \n
3.2354430379746835, \n            4.0\n        ], \n
\\"semantic_type\\": \"\", \n          \\"description\\": \"\\n      \"}, \n
\\"column\\": \\"goout\\\", \n        \\"properties\\": {\n          \\"dtype\\": \\"number\\\", \n          \\"std\\": 138.68948196584594, \n
\\"min\\": 1.0, \n          \\"max\\": 395.0, \n
\\"num_unique_values\\": 8, \n          \\"samples\\": [\n            395.0, \n
3.108860759493671, \n            3.0\n        ], \n
\\"semantic_type\\": \"\", \n          \\"description\\": \"\\n      \"}, \n
\\"column\\": \\"Dalc\\\", \n        \\"properties\\": {\n          \\"dtype\\": \\"number\\\", \n          \\"std\\": 139.0354623650101, \n
\\"min\\": 0.8907414280909659, \n          \\"max\\": 395.0, \n
\\"num_unique_values\\": 6, \n          \\"samples\\": [\n            395.0, \n
1.481012658227848, \n            5.0\n        ], \n
\\"semantic_type\\": \"\", \n          \\"description\\": \"\\n      \"}, \n
\\"column\\": \\"Walc\\\", \n        \\"properties\\": {\n          \\"dtype\\": \\"number\\\", \n          \\"std\\": 138.87302263653973, \n
\\"min\\": 1.0, \n          \\"max\\": 395.0, \n
\\"num_unique_values\\": 7, \n          \\"samples\\": [\n            395.0, \n
2.2911392405063293, \n            3.0\n        ], \n
\\"semantic_type\\": \"\", \n          \\"description\\": \"\\n      \"}, \n
\\"column\\": \\"health\\\", \n        \\"properties\\": {\n          \\"dtype\\": \\"number\\\", \n          \\"std\\": \n\"
138.50262599778412, \n          \\"min\\": 1.0, \n          \\"max\\": 395.0, \n
\\"num_unique_values\\": 7, \n          \\"samples\\": [\n            395.0, \n
3.5544303797468353, \n            4.0\n        ], \n
\\"semantic_type\\": \"\", \n          \\"description\\": \"\\n      \"}, \n
\\"column\\": \\"absences\\\", \n        \\"properties\\": {\n          \\"dtype\\": \\"number\\\", \n          \\"std\\": \n\"
136.85777166785417, \n          \\"min\\": 0.0, \n          \\"max\\": 395.0, \n
\\"num_unique_values\\": 7, \n          \\"samples\\": [\n            395.0, \n
5.708860759493671, \n            8.0\n        ], \n
\\"semantic_type\\": \"\", \n          \\"description\\": \"\\n      \"}, \n
\\"column\\": \\"G1\\\", \n        \\"properties\\": {\n          \\"dtype\\": \\"number\\\", \n          \\"std\\": 136.30663508587594, \n
\\"min\\": 3.0, \n          \\"max\\": 395.0, \n
\\"num_unique_values\\": 8, \n          \\"samples\\": [\n            10.90886075949367, \n
11.0, \n            395.0\n        ], \n
\\"semantic_type\\": \"\", \n          \\"description\\": \"\\n      \"}, \n
\\"column\\": \\"G2\\\", \n        \\"properties\\": {\n          \\"dtype\\": \\"number\\\", \n          \\"min\\": 0.0, \n          \\"max\\": \n\"

```

```

395.0,\n          "num_unique_values": 8,\n          "samples": [\n10.713924050632912,\n              11.0,\n                  395.0\n          ],\n          "semantic_type": "\",\n              "description": \"\"\n          }\n      },\n      {\n          "column": "G3",\n          "properties": {\n              "dtype": "number",\n                  "std": 136.35024783099098,\n              "min": 0.0,\n                  "max": 395.0,\n                  "num_unique_values": 8,\n          "samples": [\n              10.415189873417722,\n11.0,\n                  395.0\n          ],\n          "semantic_type": "\",\n          "description": \"\"\n      }\n  }\n},\n{"type": "dataframe"}\n\n# Converting all the non-numerical columns in dataset to numerical\n# columns using Label Encoder\nfrom sklearn.preprocessing import LabelEncoder\nlabel = LabelEncoder()\nlabel.fit_transform(student_per)\nencode_colm = student_per.iloc[:,\n[0,1,3,4,5,8,9,10,11,15,16,17,18,19,20,21,22]]\nfor i in encode_colm:\n    student_per[i] = label.fit_transform(student_per[i])\nstudent_per\n\n{"type": "dataframe", "variable_name": "student_per"}\n\n# Correlation Matrix\n\n\ncorrelation_Matrix = student_per.corr()\ncorrelation_Matrix\n\n{"type": "dataframe", "variable_name": "correlation_Matrix"}\n\n# Scatter Matrix\n\npd.plotting.scatter_matrix(student_per, figsize=(15, 10))\nplt.show()

```



```

# Storing the input attributes and the target variable
X = student_per.iloc[:, :32] # Input attributes
y = student_per.iloc[:, 32] # Target variable(G3)
X

{"type": "dataframe", "variable_name": "X"}

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test =
train_test_split(X, y, test_size=0.20, random_state=42)

from sklearn.preprocessing import StandardScaler
stand_scale = StandardScaler()
stand_scale.fit(X_train)
X_train_std = stand_scale.transform(X_train)
X_test_std = stand_scale.transform(X_test)

from sklearn.linear_model import LinearRegression
LR_model = LinearRegression()
LR_model.fit(X_train_std, y_train)
LR_score = LR_model.score(X_test_std, y_test) # Finding the R2 score of
# the model
print("R2 score:", LR_score)

R2 score: 0.7545777855043497

```

```
#Hyper parameter tuning
```

```
#1. SVM with RBF kernel.
```

```
from sklearn.svm import SVR
rbf_model = SVR(kernel='rbf')

rbf_model.get_params()

{'C': 1.0,
 'cache_size': 200,
 'coef0': 0.0,
 'degree': 3,
 'epsilon': 0.1,
 'gamma': 'scale',
 'kernel': 'rbf',
 'max_iter': -1,
 'shrinking': True,
 'tol': 0.001,
 'verbose': False}
```

Tuning the hyperparameters C and gamma

```
hyp_grid_rbf = [
    'C': [0.01, 0.1, 1, 10, 100, 100],
    'gamma': [0.001, 0.009, 1]
]
hyp_grid_rbf

[{'C': [0.01, 0.1, 1, 10, 100, 100], 'gamma': [0.001, 0.009, 1]}]
```

Performing the Grid Search

```
from sklearn.model_selection import GridSearchCV
grid_search_rbf =
GridSearchCV(estimator=rbf_model, param_grid=hyp_grid_rbf, cv=5)
grid_search_rbf.fit(X_train_std, y_train)
print("Best Parameters:", grid_search_rbf.best_params_)
print("Best Score:", grid_search_rbf.best_score_)
print("Best Estimator:", grid_search_rbf.best_estimator_)

Best Parameters: {'C': 100, 'gamma': 0.001}
Best Score: 0.8230913992113609
Best Estimator: SVR(C=100, gamma=0.001)
```

R2 score of the model after the hyperparameter tuning

```
from sklearn.svm import SVR
rbf_model_hyp = SVR(kernel='rbf', C=100, gamma=0.001)
rbf_model_hyp.fit(X_train_std, y_train)
```

```
rbf_score = rbf_model_hyp.score(X_test_std,y_test)
print("R2 Score:",rbf_score)
```

```
R2 Score: 0.7851580302181262
```

#2. Random Forest

```
from sklearn.ensemble import RandomForestRegressor
RF_model = RandomForestRegressor(random_state=42)

RF_model.get_params()

{'bootstrap': True,
 'ccp_alpha': 0.0,
 'criterion': 'squared_error',
 'max_depth': None,
 'max_features': 1.0,
 'max_leaf_nodes': None,
 'max_samples': None,
 'min_impurity_decrease': 0.0,
 'min_samples_leaf': 1,
 'min_samples_split': 2,
 'min_weight_fraction_leaf': 0.0,
 'monotonic_cst': None,
 'n_estimators': 100,
 'n_jobs': None,
 'oob_score': False,
 'random_state': 42,
 'verbose': 0,
 'warm_start': False}
```

We are tuning the

hyperparameters(n_estimators,maximum_depth,maximum_leaf_nodes,maximum_features,mi
nimum_samples_split and minimum_samples_leaf)

```
hyp_grid_rf = {'n_estimators':[i for i in range(10,101,10)],
               'max_depth':[j for j in range(2,21,2)],
               'max_leaf_nodes':[k for k in range(20,41,2)],
               'max_features':[l for l in range(2,33,2)],
               'min_samples_split':[m for m in range(2,21,2)],
               'min_samples_leaf':[n for n in range(1,21,2)]}

hyp_grid_rf

{'n_estimators': [10, 20, 30, 40, 50, 60, 70, 80, 90, 100],
 'max_depth': [2, 4, 6, 8, 10, 12, 14, 16, 18, 20],
 'max_leaf_nodes': [20, 22, 24, 26, 28, 30, 32, 34, 36, 38, 40],
 'max_features': [2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28,
 30, 32],
```

```
'min_samples_split': [2, 4, 6, 8, 10, 12, 14, 16, 18, 20],  
'min_samples_leaf': [1, 3, 5, 7, 9, 11, 13, 15, 17, 19]}
```

Performing Random Search

```
from sklearn.model_selection import RandomizedSearchCV  
random_search_rf =  
RandomizedSearchCV(estimator=RF_model,param_distributions=hyp_grid_rf,  
cv=5)  
random_search_rf.fit(X_train_std,y_train)  
print("Best Parameters:",random_search_rf.best_params_)  
print("Best Score:",random_search_rf.best_score_)  
print("Best Estimator:",random_search_rf.best_estimator_)  
  
Best Parameters: {'n_estimators': 10, 'min_samples_split': 6,  
'min_samples_leaf': 11, 'max_leaf_nodes': 36, 'max_features': 32,  
'max_depth': 12}  
Best Score: 0.8589145158962216  
Best Estimator: RandomForestRegressor(max_depth=12, max_features=32,  
max_leaf_nodes=36,  
min_samples_leaf=11, min_samples_split=6,  
n_estimators=10,  
random_state=42)
```

R2 Score of the model after hyperparameter tuning

```
from sklearn.ensemble import RandomForestRegressor  
RF_model_hyp =  
RandomForestRegressor(n_estimators=60,max_depth=18,max_leaf_nodes=22,m  
ax_features=24,min_samples_split=2,min_samples_leaf=5,n_jobs=-  
1,random_state=42)  
RF_model_hyp.fit(X_train_std,y_train)  
RF_score = RF_model_hyp.score(X_test_std,y_test)  
print("R2 Score:",RF_score)  
  
R2 Score: 0.8315030232671886
```

#3. Adaboost using decision tree stumps

```
from sklearn.ensemble import AdaBoostRegressor  
ADA_model = AdaBoostRegressor(random_state=42)  
  
ADA_model.get_params()  
  
{'estimator': None,  
'learning_rate': 1.0,  
'loss': 'linear',  
'n_estimators': 50,  
'random_state': 42}
```

We are tuning the hyperparameters n_estimators and learning rate

```
hyp_grid_ada = [{n_estimators:[i for i in range(10,201,20)],  
                 'learning_rate':np.linspace(0.15,0.5,8)}]  
hyp_grid_ada  
[{'n_estimators': [10, 30, 50, 70, 90, 110, 130, 150, 170, 190],  
 'learning_rate': array([0.15, 0.2 , 0.25, 0.3 , 0.35, 0.4 , 0.45,  
 0.5 ])}]
```

Performing the Grid Search

```
from sklearn.model_selection import GridSearchCV  
grid_search_ada =  
GridSearchCV(estimator=ADA_model,param_grid=hyp_grid_ada,cv=5)  
grid_search_ada.fit(X_train_std,y_train)  
print("The Best Parameters:",grid_search_ada.best_params_)  
print("The Best Score:",grid_search_ada.best_score_)  
print("The Best Estimator:",grid_search_ada.best_estimator_)  
  
The Best Parameters: {'learning_rate':  
np.float64(0.4499999999999996), 'n_estimators': 70}  
The Best Score: 0.8968550900297902  
The Best Estimator:  
AdaBoostRegressor(learning_rate=np.float64(0.4499999999999996),  
                  n_estimators=70, random_state=42)
```

R2 Score of the model after hyperparameter tuning

```
from sklearn.ensemble import AdaBoostRegressor  
ADA_model_hyp =  
AdaBoostRegressor(n_estimators=70,learning_rate=0.25,random_state=42)  
ADA_model_hyp.fit(X_train_std,y_train)  
ADA_score = ADA_model_hyp.score(X_test_std,y_test)  
print("R2 Score:",ADA_score)  
  
R2 Score: 0.8325293553061772
```

#4. Gradient Boosting

```
from sklearn.ensemble import GradientBoostingRegressor  
GBR_model = GradientBoostingRegressor(random_state=42)  
  
GBR_model.get_params()  
  
{'alpha': 0.9,  
 'ccp_alpha': 0.0,  
 'criterion': 'friedman_mse',  
 'init': None,  
 'learning_rate': 0.1,
```

```

'loss': 'squared_error',
'max_depth': 3,
'max_features': None,
'max_leaf_nodes': None,
'min_impurity_decrease': 0.0,
'min_samples_leaf': 1,
'min_samples_split': 2,
'min_weight_fraction_leaf': 0.0,
'n_estimators': 100,
'n_iter_no_change': None,
'random_state': 42,
'subsample': 1.0,
'tol': 0.0001,
'verbose': 0,
'warm_start': False}

```

We are tuning the hyperparameters n_estimators,max_depth and learning_rate

```

hyp_grid_gbr = {'n_estimators':[i for i in range(10,101,10)],
                'learning_rate':np.linspace(0.1,0.5,5),
                'max_depth':[j for j in range(2,21,2)]}
hyp_grid_gbr
{'n_estimators': [10, 20, 30, 40, 50, 60, 70, 80, 90, 100],
 'learning_rate': array([0.1, 0.2, 0.3, 0.4, 0.5]),
 'max_depth': [2, 4, 6, 8, 10, 12, 14, 16, 18, 20]}

```

Performing Random Search

```

from sklearn.model_selection import RandomizedSearchCV
random_search_gbr =
RandomizedSearchCV(estimator=GBR_model,param_distributions=hyp_grid_gb
r, cv=5)
random_search_gbr.fit(X_train_std,y_train)
print("Best Parameters:",random_search_gbr.best_params_)
print("Best Score:",random_search_gbr.best_score_)
print("Best Estimator:",random_search_gbr.best_estimator_)

Best Parameters: {'n_estimators': 10, 'max_depth': 4, 'learning_rate': np.float64(0.30000000000000004)}
Best Score: 0.8867884616886164
Best Estimator:
GradientBoostingRegressor(learning_rate=np.float64(0.30000000000000004),
),
                                         max_depth=4, n_estimators=10,
random_state=42)

```

R2 Score of the model after hyperparamter tuning

```

from sklearn.ensemble import GradientBoostingRegressor
GBR_model_hyp =
GradientBoostingRegressor(n_estimators=20,learning_rate=0.3,max_depth=4,random_state=42)
GBR_model_hyp.fit(X_train_std,y_train)
GBR_score = GBR_model_hyp.score(X_test_std,y_test)
print("R2 Score:",GBR_score)

R2 Score: 0.811762031831491

```

Decision Tree

```

from sklearn.tree import DecisionTreeRegressor
tree_model = DecisionTreeRegressor(random_state=42)
tree_model.get_params()

{'ccp_alpha': 0.0,
 'criterion': 'squared_error',
 'max_depth': None,
 'max_features': None,
 'max_leaf_nodes': None,
 'min_impurity_decrease': 0.0,
 'min_samples_leaf': 1,
 'min_samples_split': 2,
 'min_weight_fraction_leaf': 0.0,
 'monotonic_cst': None,
 'random_state': 42,
 'splitter': 'best'}

```

We are tuning the hyperparameters

max_depth,max_features,max_leaf_nodes,min_samples_split and min_samples_leaf

```

hyp_grid_tree = {'max_depth':[i for i in range(2,21,2)],
                 'max_features':[j for j in range(2,33,2)],
                 'max_leaf_nodes':[k for k in range(20,41,2)],
                 'min_samples_split':[l for l in range(2,21,2)],
                 'min_samples_leaf':[m for m in range(1,21,2)]}
hyp_grid_tree

{'max_depth': [2, 4, 6, 8, 10, 12, 14, 16, 18, 20],
 'max_features': [2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32],
 'max_leaf_nodes': [20, 22, 24, 26, 28, 30, 32, 34, 36, 38, 40],
 'min_samples_split': [2, 4, 6, 8, 10, 12, 14, 16, 18, 20],
 'min_samples_leaf': [1, 3, 5, 7, 9, 11, 13, 15, 17, 19]}

```

Performing Random Search

```

from sklearn.model_selection import RandomizedSearchCV
random_search_tree =
RandomizedSearchCV(estimator=tree_model,param_distributions=hyp_grid_t
ree,cv=5)
random_search_tree.fit(X_train_std,y_train)
print("Best Parameters:",random_search_tree.best_params_)
print("Best Score:",random_search_tree.best_score_)
print("Best Estimator:",random_search_tree.best_estimator_)

Best Parameters: {'min_samples_split': 4, 'min_samples_leaf': 7,
'max_leaf_nodes': 22, 'max_features': 24, 'max_depth': 4}
Best Score: 0.7806317573935685
Best Estimator: DecisionTreeRegressor(max_depth=4, max_features=24,
max_leaf_nodes=22,
                                         min_samples_leaf=7, min_samples_split=4,
random_state=42)

```

R2 score of the model after hyperparameter tuning

```

from sklearn.tree import DecisionTreeRegressor
tree_model_hyp =
DecisionTreeRegressor(max_depth=8,max_features=28,max_leaf_nodes=20,mi
n_samples_split=8,min_samples_leaf=7,random_state=42)
tree_model_hyp.fit(X_train_std,y_train)
tree_score = tree_model_hyp.score(X_test_std,y_test)
print("R2 Score:",tree_score)

R2 Score: 0.8493239751834084

import matplotlib.pyplot as plt

# Example model names – adjust if you know the actual ones
model_names = [
    "Linear Regression", "Random Search SVM RBF", "Random Search
Random Forest", "Random Search Ada Boost",
    "Random Search Gradient Boost", "Random Search Decision Tree"
]
r2_scores = [ LR_score, rbf_score, RF_score, ADA_score, GBR_score,
tree_score ]

# Create the bar chart
plt.figure(figsize=(12, 6))
bars = plt.bar(model_names, r2_scores, color='cornflowerblue')

# Add R2 values on top of each bar
for bar in bars:
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, yval + 0.01,
f'{yval:.2f}', ha='center', va='bottom')

```

```
# Labeling
plt.title("Model Performance (R2 Score)")
plt.xlabel("Model Type")
plt.ylabel("R2 Score")
plt.ylim(0, 1.1)
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```

